

ASSESSMENT TOOL 3 — Annexure A: Integrated Experiments

Name: B.Satya Hyndavi

Register Number: 192311178

Course Code: CSA1704

Experiment 1: Decision Tree Classification (Iris dataset)

(a) Load, pre-process, split

Dataset: the Iris dataset (150 samples, 4 numeric features: sepal length/width, petal length/width; target = species, 3 classes).

	sepal length	sepal width	petal length	petal width	target	species
0	5.1	3.5	1.4	0.2	0	setosa
1	4.9	3.0	1.4	0.2	0	setosa
2	4.7	3.2	1.3	0.2	0	setosa
3	4.6	3.1	1.5	0.2	0	setosa
4	5.0	3.6	1.4	0.2	0	setosa

Data types: all 4 features are float64; target is int64 (already label-encoded 0/1/2 by scikitlearn); species is the human-readable string label. Missing-value check: 0 missing values in every column — no imputation needed for this dataset.

Preprocessing note: Iris ships fully numeric with no missing data, so the encoding/imputation steps requested by the brief have nothing to act on here; if a categorical feature or missing values were present (as in a Play-Tennis-style dataset), the same pipeline would add `pandas.get_dummies()/LabelEncoder` for categoricals and `SimpleImputer` for missing values before the split below.

Split: 70% train / 30% test, stratified by class → Train = 105 rows, Test = 45 rows.

(b) Build the Decision Tree and justify the root node




Fig. 1 — Trained Decision Tree (Gini criterion, max_depth=4)

Root split: petal length (cm) ≤ 2.45

Root Gini impurity = 0.6667 $\rightarrow$ Left child Gini = 0.0000 Right child Gini = 0.5000

Justification: the root is chosen by the Gini index (equivalently, Information-Gain-style impurity reduction) — among all 4 features and all possible thresholds, splitting on petal length ≤ 2.45 gives the single largest impurity drop: it perfectly isolates all 35 setosa samples into a pure leaf (Gini = 0, 100% setosa) in one step, cutting the parent's Gini from 0.667 to an average child Gini of just 0.25. Feature importances computed by the fitted tree confirm this numerically: petal length = 0.549 (highest), petal width = 0.436, sepal width = 0.014, sepal length = 0.000 — petal length is by far the most discriminative feature, which is exactly why the algorithm selects it as the root.

(c) Evaluation

Metric	Value
Accuracy	0.8889 (40/45 correct)

Confusion matrix (rows = actual, columns = predicted; order: setosa, versicolor, virginica):

	Pred: setosa	Pred: versicolor	Pred: virginica
Actual: setosa	15	0	0
Actual: versicolor	0	12	3
Actual: virginica	0	2	13

Class	Precision	Recall	F1-score	Support
setosa	1.00	1.00	1.00	15
versicolor	0.86	0.80	0.83	15
virginica	0.81	0.87	0.84	15

Comment on performance: setosa is separated perfectly (linearly separable from the other two species on petal measurements alone, matching the pure root split above). All 5 errors are versicolor ↔ virginica confusions — these two species genuinely overlap in petal length/width, so this is expected, benign error rather than a modelling flaw; a shallow, 4-level tree still resolves 40/45 test cases correctly.

At least two misclassified instances (from the actual test run):

- Row (petal length=4.7, petal width=1.6, sepal length=6.3, sepal width=3.3): true = versicolor, predicted = virginica.
- Row (petal length=4.8, petal width=1.8, sepal length=6.0, sepal width=3.0): true = virginica, predicted = versicolor.
- (3 further versicolor/virginica confusions occurred in the same borderline petal-width region, ≈1.6–1.7 cm — right at the tree's 1.55/1.70 cm decision thresholds.)

Code (executed with Python / scikit-learn)

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```

iris = load_iris(as_frame=True) df
= iris.frame.copy()
X, y = df[iris.feature_names], df['target']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42, stratify=y)

dt = DecisionTreeClassifier(criterion="gini", max_depth=4, random_state=42)
dt.fit(X_train, y_train) print(export_text(dt,
feature_names=list(iris.feature_names)))

y_pred = dt.predict(X_test) print("Accuracy:", accuracy_score(y_test,
y_pred)) print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred, target_names=iris.target_names))

```

Experiment 2: Reinforcement Learning — Q-Learning on 4×4 Grid World

(a) Environment definition

- States: 16 grid cells, indexed (row, col), $\text{row/col} \in \{0,1,2,3\}$.
- Actions: {Up, Down, Left, Right} — moving off the grid edge keeps the agent in place.
- Start = (0,0); Goal = (3,3); Obstacle = (1,1).
- Reward structure: +10 for reaching the Goal, −1 for every ordinary step, −5 for bumping into the Obstacle (agent stays in place).
- Q-table initialised to all zeros, shape 16 states $\times$ 4 actions.
- Hyperparameters: learning rate $\alpha = 0.1$, discount factor $\gamma = 0.9$, exploration rate $\epsilon = 0.2$ (ϵ -greedy).

(b) Q-Learning run — 100 episodes

Q-values for 2 representative states at episodes 1, 50 and 100 (order of the 4 numbers = Up, Down, Left, Right):

State (0,0) — the START state:

Episode 1: [-0.199, -0.190, -0.199, -0.190]

Episode 50: [-2.283, -2.242, -2.267, -2.255]

Episode 100: [-2.195, 0.693, -2.151, -2.301] -> best action settles on Down

State (2,2) — a mid-grid state near the goal:

Episode 1: [-0.100, 0.000, -0.100, -0.100]

Episode 50: [-0.408, 0.151, -0.208, 0.811]

Episode 100: [-0.335, 0.151, -0.208, 3.891] -> best action settles on Right

How the Q-table evolves: at episode 1 the values are tiny (only 1 step of update has touched each entry, mostly reflecting the -1 step cost). By episode 50 the state (0,0) values are still all deeply negative and close together — early in training the agent is still bumping into the obstacle and wandering under ϵ -greedy exploration, so no action has pulled clearly ahead yet. By episode 100 a clear winner has emerged at each tracked state (Down at (0,0), Right at (2,2)), and Q-values near the goal (e.g. (2,2)→Right = 3.89) have grown substantially larger than Q-values further away (e.g. (0,0)→Down = 0.69) — exactly the expected shape of a converging Q-function, where value propagates backward from the +10 goal reward toward the start state as episodes accumulate.

Cumulative reward per episode



Fig. 2 — Cumulative reward per episode, with 5-episode moving average

Convergence behaviour: the first 10 episodes average a cumulative reward of -17.90 per episode — heavy exploration, frequent obstacle bumps, and long, inefficient paths to the goal (or timing out at 50 steps without reaching it). By the last 10 episodes the average has risen to $+3.20$ per episode, and the moving-average curve flattens out from roughly episode 25 onward, oscillating in a narrow positive band. This is justified by the Q-learning convergence theory: as more (state, action) pairs are visited and their estimates propagate backward from the terminal $+10$ reward via the Bellman update, the greedy policy stabilises; the remaining episode-to-episode variance is expected and comes from the $\epsilon = 0.2$ exploration rate never fully switching off, so roughly 1 in 5 actions remains randomly chosen even after the policy has effectively converged.

(c) Final learned policy




Fig. 3 — Final greedy policy at every state (arrows), obstacle and goal marked

The learned policy routes the agent from every reachable state toward the goal at (3,3) while steering around the obstacle at (1,1) — e.g. from the start (0,0) the policy first moves Down to (1,0), then continues Down/Right along the left and bottom edges, deliberately avoiding column 1/row 1's obstacle cell rather than cutting through it, which is exactly the behaviour the -5 obstacle penalty was designed to produce. **Code (executed with Python / NumPy)**

```
import numpy as np
```

```
N = 4; ACTIONS = ["Up", "Down", "Left", "Right"]
```

```
GOAL, OBSTACLE, START = (3,3), (1,1), (0,0)
```

```

def step(s, a):    r, c
= divmod(s, N)    nr,
nc = r, c
    if ACTIONS[a]=="Up": nr = max(0, r-1)    elif
ACTIONS[a]=="Down": nr = min(N-1, r+1)
elif ACTIONS[a]=="Left": nc = max(0, c-1)    elif
ACTIONS[a]=="Right": nc = min(N-1, c+1)    if
(nr,nc) == OBSTACLE: return s, -5, False
s_next = nr*N + nc    if (nr,nc) == GOAL: return
s_next, 10, True    return s_next, -1, False

alpha, gamma, epsilon = 0.1, 0.9, 0.2
Q = np.zeros((N*N, 4))

for ep in range(1, 101):
s = 0 # start state (0,0)
for t in range(50):
    a = np.random.randint(4) if np.random.rand()<epsilon else int(np.argmax(Q[s]))
s_next, r, done = step(s, a)
    Q[s,a] += alpha * (r + gamma*np.max(Q[s_next]) - Q[s,a])
s = s_next    if done: break

policy = {divmod(s,N): ACTIONS[int(np.argmax(Q[s]))] for s in range(N*N)}

```